



An on-call QA engineer for your AI agents. Read the spoken lines; let the board do the rest.

0:00 Cold open

STAGE: board on screen, nothing clicked.

"Raise your hand if you've shipped an AI agent this year. *(beat)* Keep it up if you'd bet **your job** that not one of them has quietly gotten **worse** since launch. *(beat — hands drop)*

That's the problem. An agent doesn't crash when it regresses. It just starts being **wrong** — confidently, in production — **and your tests stay green**.

Self-healing a renamed button? Solved. Commodity. Catching the moment your agent's **judgment** slips? **Nobody is governing that**.

So we built **TestPilot — an on-call QA engineer for your AI agents**. This is Mission Control. Every node is a live agent under continuous evaluation. Fleet trust: **seventy-four**. And this number — **eighteen bad merges prevented** — is the one I want you to watch."

Recording alt-open (no audience): "An AI agent doesn't crash when it regresses — it just starts being wrong, confidently, in production, while your tests stay green. **TestPilot** is the on-call QA engineer that catches it."

0:30 Triage, live

STAGE: click the pulsing ► PLAY DEMO.

"One evaluation run just went **red**. Watch TestPilot triage it in real time.

STAGE: green packet → PR lane.

A renamed selector — *btn-signin* to *btn-login*. **Mechanical drift**. It writes the one-line fix and opens a **GitHub pull request**. Safe to automate. Done.

STAGE: amber packet → Q lane.

A **flaky** test that passes on retry. **Quarantined** with a retry policy. Not a line of code touched.

STAGE: red packet approaches the firewall.

And now — the one that matters. Invoice-Dispute. Faithfulness just fell from a hundred... to **forty-one**. The agent stopped citing policy before it answered. That is **not** a bug. **The agent changed its mind.**"

1:15 The firewall — the whole product in one second

STAGE: red packet bounces off the FIREWALL → ESC; the BAD MERGES counter ticks up.

"And **this** is the entire idea. TestPilot **refuses**. It will **not** auto-fix behavior. It throws it at the **firewall** and escalates to a human — because **a behavioral change is a product decision, not a bug fix**.

Read it off the screen with me: **only a human merges behavior**. *(counter ticks)* And there's our number — **nineteen**.

STAGE: evidence drawer opens.

And the human doesn't get a stack trace. They get the **tool path**, the exact divergence, every eval score, and an **AI root-cause** — keyless, through the **UiPath AI Trust Layer**."

1:55 Tamper-evident ledger

STAGE: demo lands on the Governance Ledger; click TAMPER DEMO, then click it again.

“Every one of those decisions lands in a **hash-chained, tamper-evident ledger**. Don't take my word for it — let me **forge** an entry. *(chain turns red)* **Chain: broken**. Instantly. Revert it... *(green)* **Verified**. You cannot quietly rewrite this history.”

2:20 Prove the policy is real

STAGE: press Esc, then click the red “ Try to auto-merge behavior”.

“And the policy isn't a promise on a slide — it's **code**. So let me try to **cheat**. Let me approve a behavior change with my own hands... *(giant DENIED stamp slams in)* **Denied**. Not even the operator gets to merge behavior.”

2:40 Stack + close

STAGE: gesture to the UIPATH PIPELINE strip.

“All of it is pure UiPath — **Maestro** orchestration, **coded agents** on serverless, **Agent Evaluations** as the test bed, **Integration Service** for GitHub and Slack, **Action Center** for the human gate.

And it was built **with** UiPath for Coding Agents — **Claude Code** wrote these agents test-first. **Eighty-three tests. All green.**

Everybody can heal a selector. **Almost nobody can govern a behavior**. That's the moat. That's **TestPilot**. Thank you.”

30-SECOND ELEVATOR

“TestPilot is an **on-call QA engineer for AI agents**. When an evaluation goes red, it triages: mechanical drift is **auto-healed into a GitHub PR**, flaky tests are **quarantined**, and a **behavioral regression** — where the agent's judgment actually changed — is **never auto-fixed**. It hits a firewall and escalates to a human, because **only a human merges behavior**. Every decision is written to a **tamper-evident, hash-chained ledger**. Built on Maestro, coded agents, Agent Evaluations, and the AI Trust Layer — authored test-first with UiPath for Coding Agents.”

THREE LINES TO LAND

1. “**An on-call QA engineer for your AI agents.**”
2. “**A behavioral change is a product decision, not a bug fix — only a human merges behavior.**”
3. “**Everybody can heal a selector; almost nobody can govern a behavior. That's the moat.**”

“A behavioral change in an agent is a product decision, not a bug fix — **only a human merges behavior.**”